

Titanic Detailed Analysis

Kushal Kapadia

October 7, 2017

```
# Loading the new data
train = read.csv("train.csv", header = T)
test = read.csv("test.csv", header = T)

# Combining the dataset the old way by adding a column instead of doing
bind_rows fom dplyr package
test$survived = data.frame(Survived = rep("None", nrow(test)), test[,])
data.combined = rbind(train,test$survived)

# Let's have a close look at the structure of the data
str(data.combined)

## 'data.frame':    1309 obs. of  12 variables:
## $ PassengerId: int  1 2 3 4 5 6 7 8 9 10 ...
## $ Survived   : chr  "0" "1" "1" "1" ...
## $ Pclass     : int   3 1 3 1 3 3 1 3 3 2 ...
## $ Name       : Factor w/ 1307 levels "Abbing, Mr. Anthony",...: 109 191
358 277 16 559 520 629 417 581 ...
## $ Sex        : Factor w/ 2 levels "female","male": 2 1 1 1 2 2 2 2 1 1
...
## $ Age        : num   22 38 26 35 35 NA 54 2 27 14 ...
## $ SibSp      : int    1 1 0 1 0 0 0 3 0 1 ...
## $ Parch      : int    0 0 0 0 0 0 0 1 2 0 ...
## $ Ticket     : Factor w/ 929 levels "110152","110413",...: 524 597 670 50
473 276 86 396 345 133 ...
## $ Fare       : num    7.25 71.28 7.92 53.1 8.05 ...
## $ Cabin      : Factor w/ 187 levels "", "A10", "A14",...: 1 83 1 57 1 1 131
1 1 1 ...
## $ Embarked   : Factor w/ 4 levels "", "C", "Q", "S": 4 2 4 4 4 3 4 4 4 2 ...

# Letting R know the categorical variables in the dataset
data.combined$Survived = as.factor(data.combined$Survived)
data.combined$Pclass = as.factor(data.combined$Pclass)

# Let's have a general look at how many survived and how many did not
table(data.combined$Survived)


##
##      0      1 None
## 549   342   418

# Distribution across classes
table(data.combined$Pclass)
```

```
##
## 1 2 3
## 323 277 709

# Data visualization Library
library(ggplot2)

train$Pclass = as.factor(train$Pclass)
ggplot(train, aes(x = Pclass, fill = factor(Survived))) + geom_bar(width =
0.5) + xlab("Pclass") + ylab("Total Count") + labs(fill="Survived")
```



```
# Examine the first few names in the training dataset
head(as.character(train$Name))

## [1] "Braund, Mr. Owen Harris"
## [2] "Cumings, Mrs. John Bradley (Florence Briggs Thayer)"
## [3] "Heikkinen, Miss. Laina"
## [4] "Futrelle, Mrs. Jacques Heath (Lily May Peel)"
## [5] "Allen, Mr. William Henry"
## [6] "Moran, Mr. James"

# how many unique names across both train and test dataset?
length(unique(as.character(data.combined$Name)))

## [1] 1307

#Thus, this shows us that there are two duplicate names
```

```

# Finding the two duplicate names
which(duplicated(data.combined$Name))

## [1] 892 898

# One way to find the location of the two similar names
which(data.combined$Name == "Connolly, Miss. Kate")

## [1] 290 898

library(stringr) #character extraction library

# Let's try to find the observations having only "Miss." using the str_detect
misses = data.combined[which(str_detect(data.combined$Name, "Miss.")), ]
misses[1:5,]

##      PassengerId Survived Pclass                                Name      Sex
## 3              3         1      3                Heikkinen, Miss. Laina female
## 11             11         1      3      Sandstrom, Miss. Marguerite Rut female
## 12             12         1      1                Bonnell, Miss. Elizabeth female
## 15             15         0      3 Vestrom, Miss. Hulda Amanda Adolfina female
## 23             23         1      3      McGowan, Miss. Anna "Annie" female
##      Age SibSp Parch      Ticket    Fare Cabin Embarked
## 3    26     0     0 STON/O2. 3101282  7.9250      S
## 11    4     1     1      PP 9549 16.7000     G6      S
## 12   58     0     0     113783 26.5500    C103      S
## 15   14     0     0     350406  7.8542      S
## 23   15     0     0     330923  8.0292      Q

# Let's also try to find the observations having only "Mrs." using the
str_detect
mrises = data.combined[which(str_detect(data.combined$Name, "Mrs.")), ]
mrises[1:5,]

##      PassengerId Survived Pclass                                Name      Sex Age SibSp
## 2              2         1      1                Cumings, Mrs. John Bradley (Florence Briggs Thayer) female 38     1
## 4              4         1      1      Futrelle, Mrs. Jacques Heath (Lily May Peel) female 35     1
## 9              9         1      3      Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg) female 27     0
## 10             10         1      2                Nasser, Mrs. Nicholas (Adele Achem) female 14     1
## 16             16         1      2      Hewlett, Mrs. (Mary D Kingcome) female 55     0
##      Parch      Ticket    Fare Cabin Embarked
## 2         0 PC 17599 71.2833     C85      C
## 4         0  113803 53.1000    C123      S
## 9         2  347742 11.1333      S

```

```
## 10      0    237736 30.0708      C
## 16      0    248706 16.0000      S
```

check out if the pattern continues

```
males = data.combined[data.combined$Sex == "male",]
males[1:5,]
```

```
##   PassengerId Survived Pclass                Name Sex Age
## 1           1         0      3   Braund, Mr. Owen Harris male  22
## 5           5         0      3   Allen, Mr. William Henry male  35
## 6           6         0      3      Moran, Mr. James male   NA
## 7           7         0      1   McCarthy, Mr. Timothy J male  54
## 8           8         0      3 Palsson, Master. Gosta Leonard male   2
##   SibSp Parch   Ticket    Fare Cabin Embarked
## 1      1     0  A/5 21171   7.2500      S
## 5      0     0  373450   8.0500      S
## 6      0     0  330877   8.4583      Q
## 7      0     0   17463  51.8625   E46      S
## 8      3     1  349909  21.0750      S
```

Now, we're gonna add variable "title" to the dataset

Before adding the title variable, we first have to assign the title values to all the observations

Creating a utility function to help with title extraction

```
extractTitle = function(name){
  name = as.character(name)


  if(length(grep("Miss.",name))>0){
    return("Miss.")
  } else if (length(grep("Mrs.",name))>0){
    return("Mrs.")
  } else if (length(grep("Master.",name))>0){
    return("Master.")
  } else if (length(grep("Mr.",name))>0){
    return("Mr.")
  } else {
    return("Other")
  }
}

titles = NULL
for(i in 1:nrow(data.combined)){
  titles = c(titles, extractTitle(data.combined[i,"Name"]))
}
data.combined$title = as.factor(titles)
```

Data Visualization

```
par(mfrow = c(3,3))
```


```
# Before ending this, let's also try visualizing data with 3 variables
together namely survived, pclass and title
ggplot(data.combined[1:891,], aes(x = title , fill = Survived)) + #remember
this to write stat=count and position=dodge
  geom_bar(stat = 'count', position = "dodge") +
  facet_wrap(~Pclass) +
  xlab("Title") +
  ylab("Total Count") +
  labs(fill = "Survived")
```



```
# Distribution of male to female in the combined dataset
table(data.combined$Sex)

##
## female    male
##    466    843

# Now's let's visualize data a bit for pclass, sex and survived using ggplot
ggplot(data.combined[1:891,], aes(x = Sex, fill=Survived)) +
  geom_bar(width = 0.5) +
  facet_wrap(~Pclass) +
  xlab("Sex") +
  ylab("Total Counts") +
  labs(fill="Survived")
```



Enough with the sex variable now. Age and sex seem to be related to each other

Let's explore the age variable a little bit more

```
summary(data.combined$Age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##      0.17  21.00   28.00   29.88  39.00   80.00     263
```

summary(data.combined[1:891,"Age"]) #This actually indicates that there are lot of missing values for age in the training data which is not a good thing.

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##      0.42  20.12   28.00   29.70  38.00   80.00     177
```


Ways to find missing values are mean, median, mean or median of the group, imputation(basically training a model which would help us in getting the missing values) and also proxy

Let's visualize again by using age, sex, pclass and survived. Should be interesting to code it! I'm excited!

```
ggplot(data.combined[1:891,], aes(x=Age, fill=Survived)) +
  geom_histogram(binwidth = 10) +
  facet_wrap(~Sex+Pclass) +
  ggtitle("Sex+Pclass") +
  xlab("Age") +
```

```
ylab("Total Count") +
labs(fill = "Survived")
```

```
## Warning: Removed 177 rows containing non-finite values (stat_bin).
```



Master is a good proxy for male children. Here's why:

```
boys = data.combined[which(data.combined$title == "Master."),]
summary(boys$Age) #This indeed confirms that the min age is 0.33 and max age
is 14.5 which means that they are male children
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##  0.330   2.000   4.000   5.483   9.000  14.500         8
```

Let's also delve deep into Miss title which is a bit complicated and we'll see why:

```
misses = data.combined[data.combined$title == "Miss.",]
summary(misses$Age) # See, actually here the min age 0.17 rises to Max 63
which is why we can't say what they exactly are. Female children or female
adults
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##  0.17   15.00   22.00   21.77   30.00   63.00        50
```

Let's visualize something different using ggplot

```
ggplot(misses[misses$Survived != "None",], aes(x=Age, fill=Survived)) +
  facet_wrap(~Pclass) +
  geom_histogram(binwidth = 5) +
  ggtitle("Age for Miss. by Pclass") +
```

```
xlab("Age") +
ylab("Total Count")
```

```
## Warning: Removed 36 rows containing non-finite values (stat_bin).
```



```
# Okay appears that female children might have a different survival rate
# Could be a candidate for feature engineering
```

```
misses.alone = misses[misses$SibSp == 0 & misses$Parch == 0,]
summary(misses.alone$Age)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
##      5.00   21.00   26.00   27.23   32.50   58.00     33
```

```
length(which(misses.alone$Age <= 14.5))
```

```
## [1] 4
```

```
# Now, Let's take a look at Sibsp variable
summary(data.combined$SibSp)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      0.0000  0.0000  0.0000  0.4989  1.0000  8.0000
```

```
table(data.combined$SibSp) #Not there in the video
```

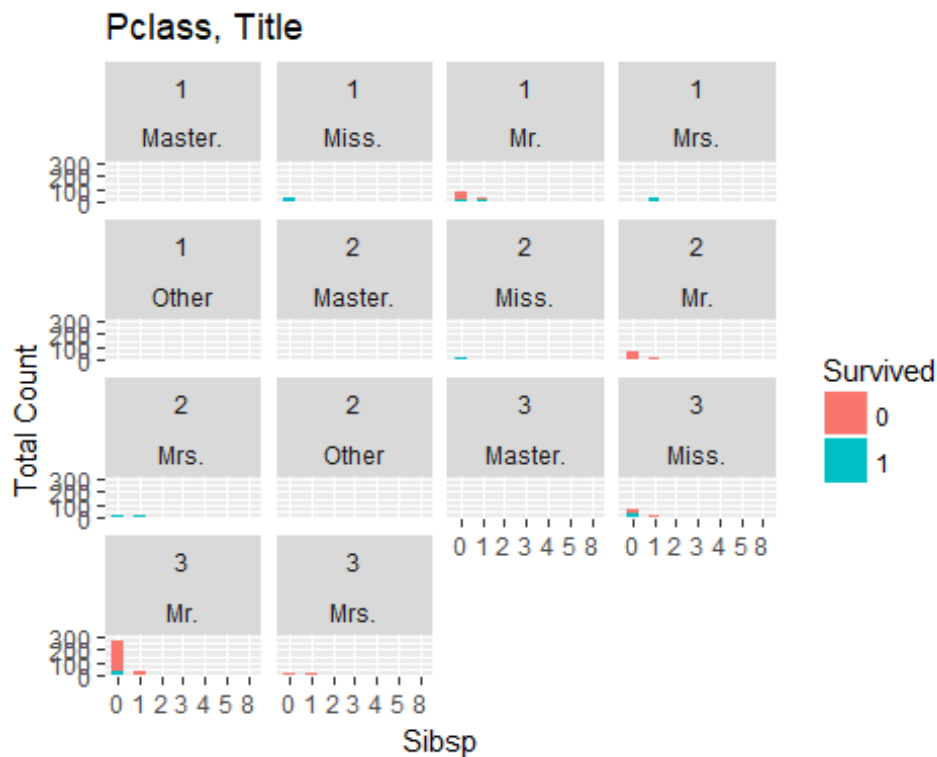
```
##
##      0      1      2      3      4      5      8
## 891 319   42   20   22    6    9
```


As very evident from the 2 above lines of code, we can turn Sibsp into a factor

```
data.combined$SibSp = as.factor(data.combined$SibSp)
```

Now that we know something about Sibsp, let's try visualizing it for a wee bit:

```
ggplot(data.combined[1:891,], aes(x = SibSp, fill = Survived)) +
  geom_bar(width = 0.5) +
  facet_wrap(~Pclass + title) +
  ggtitle("Pclass, Title") +
  xlab("Sibsp") +
  ylab("Total Count") +
  ylim(0,300) +
  labs(fill = "Survived")
```



Let's have a look at parch variable which actually means parents or something. Ain't sure enough :P

```
unique(data.combined$Parch)
```

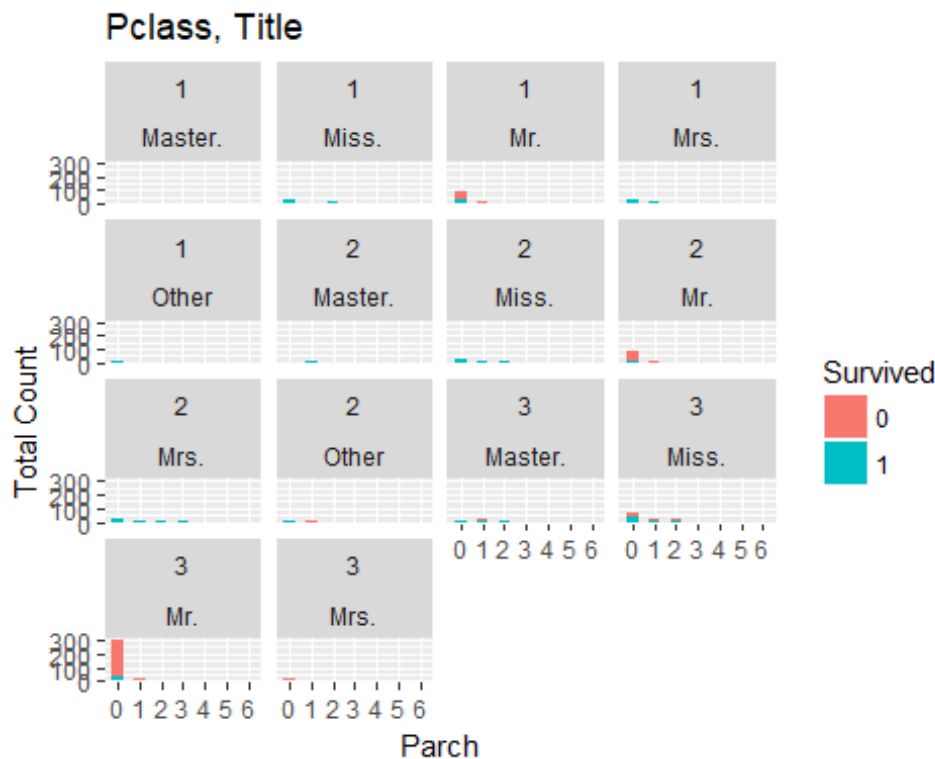
```
## [1] 0 1 2 5 3 4 6 9
```

```
table(data.combined$Parch)
```

```
##
##    0    1    2    3    4    5    6    9
## 1002 170 113    8    6    6    2    2
```

```
# Let's go ahead and convert this also to a factor variable
data.combined$Parch = as.factor(data.combined$Parch)

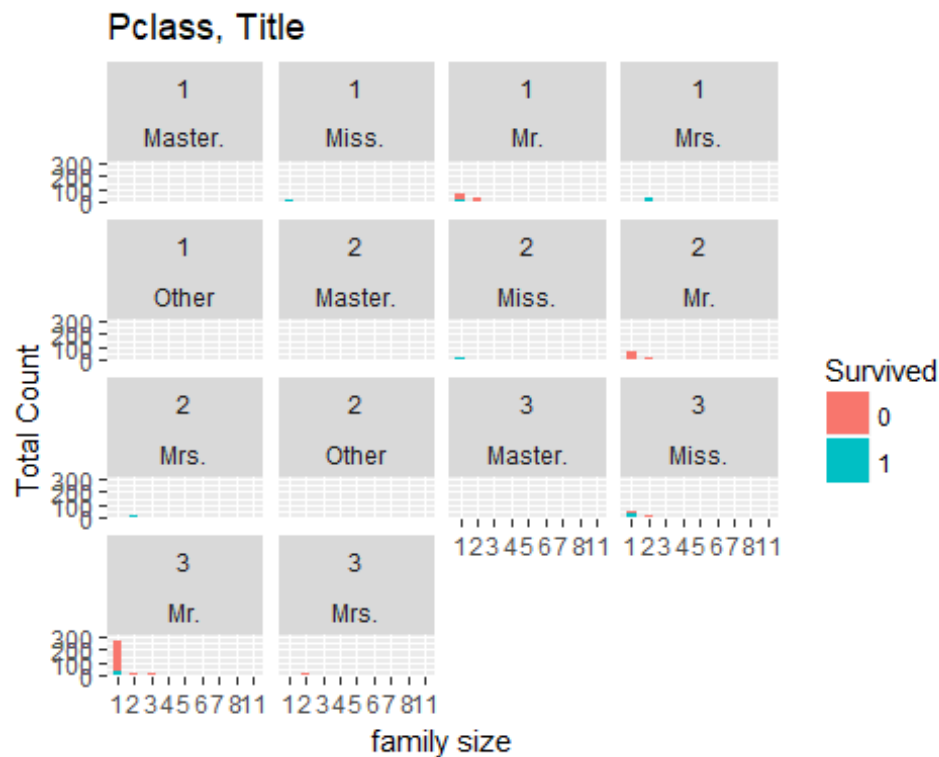
# Same kind of visualization plot as Sibsp. Hence, Copying and pasting it:
ggplot(data.combined[1:891,], aes(x = Parch, fill = Survived)) +
  geom_bar(width = 0.5) +
  facet_wrap(~Pclass + title) +
  ggtitle("Pclass, Title") +
  xlab("Parch") +
  ylab("Total Count") +
  ylim(0,300) +
  labs(fill = "Survived")
```



```
# Let's do some cool feature engineering here by creating a variable
familysize
# Let's first combine Sibsp and Parch from training and testing dataset
temp.Sibsp = c(train$SibSp, test$SibSp)
temp.Parch = c(train$Parch, test$Parch)
data.combined$familysize = as.factor(temp.Parch + temp.Sibsp + 1)

# Visualize it to see it has some predictive power in there or not:
ggplot(data.combined[1:891,], aes(x = familysize, fill = Survived)) +
  geom_bar(width = 0.5) +
  facet_wrap(~Pclass + title) +
  ggtitle("Pclass, Title") +
  xlab("family size") +
  ylab("Total Count") +
```

```
ylim(0,300) +
  labs(fill = "Survived")
```



```
par(mfrow = c(3,3))
```

We need to Look at the ticket variable

```
str(data.combined$Ticket)
```

```
## Factor w/ 929 levels "110152","110413",...: 524 597 670 50 473 276 86 396 345 133 ...
```

```
summary(data.combined$Ticket)
```

```
##          CA. 2343          1601          CA 2144
##              11              8              8
##          3101295          347077          347082
##              7              7              7
##          PC 17608          S.O.C. 14879          113781
##              7              7              6
##          19950          347088          382652
##              6              6              6
##          113503          16966          220845
##              5              5              5
##          349909          4133          PC 17757
##              5              5              5
##          W./C. 6608          113760          12749
```

##	5	4	4
##	17421	230136	24160
##	4	4	4
##	2666	36928	C.A. 2315
##	4	4	4
##	C.A. 33112	C.A. 34651	LINE
##	4	4	4
##	PC 17483	PC 17755	PC 17760
##	4	4	4
##	SC/Paris 2123	W./C. 6607	110152
##	4	4	3
##	110413	11767	13502
##	3	3	3
##	19877	19928	230080
##	3	3	3
##	239853	248727	248738
##	3	3	3
##	26360	2650	2653
##	3	3	3
##	2661	2662	2668
##	3	3	3
##	2678	28220	29103
##	3	3	3
##	29106	29750	315153
##	3	3	3
##	33638	345773	347080
##	3	3	3
##	347742	35273	363291
##	3	3	3
##	367226	370129	371110
##	3	3	3
##	A/4 48871	A/5. 851	C 4001
##	3	3	3
##	C.A. 2673	C.A. 31921	C.A. 37671
##	3	3	3
##	F.C.C. 13529	PC 17558	PC 17569
##	3	3	3
##	PC 17572	PC 17582	PC 17756
##	3	3	3
##	PC 17758	PC 17761	PP 9549
##	3	3	3
##	S.C./PARIS 2079	C.A. 31029	SOTON/O.Q. 3101315
##	3	3	3
##	110465	110813	111361
##	2	2	2
##	112058	113059	113505
##	2	2	2
##	113509	113572	113773
##	2	2	2
##	113776	113789	113796

```
##           2           2           2
##       113798       113803       113806
##           2           2           2
##       (Other)
##           947
```

Thus, we're gonna transform this into a character variable

```
data.combined$Ticket = as.character(data.combined$Ticket)
data.combined$Ticket[1:20]
```

```
## [1] "A/5 21171"      "PC 17599"      "STON/02. 3101282"
## [4] "113803"         "373450"        "330877"
## [7] "17463"          "349909"        "347742"
## [10] "237736"         "PP 9549"       "113783"
## [13] "A/5. 2151"      "347082"        "350406"
## [16] "248706"         "382652"        "244373"
## [19] "345763"         "2649"
```

Looking at a first few values, we can extract the first string to check if something useful comes out

```
ticket.first.char = ifelse(data.combined$Ticket == "", " ",
  substr(data.combined$Ticket,1,1))
unique(ticket.first.char)
```


```
## [1] "A" "P" "S" "1" "3" "2" "C" "7" "W" "4" "F" "L" "9" "6" "5" "8"
```

Okay so we can make it a factor for analysis purposes and visualize it


```
data.combined$ticket.first.char = as.factor(ticket.first.char)
```

First, a high level plot of data

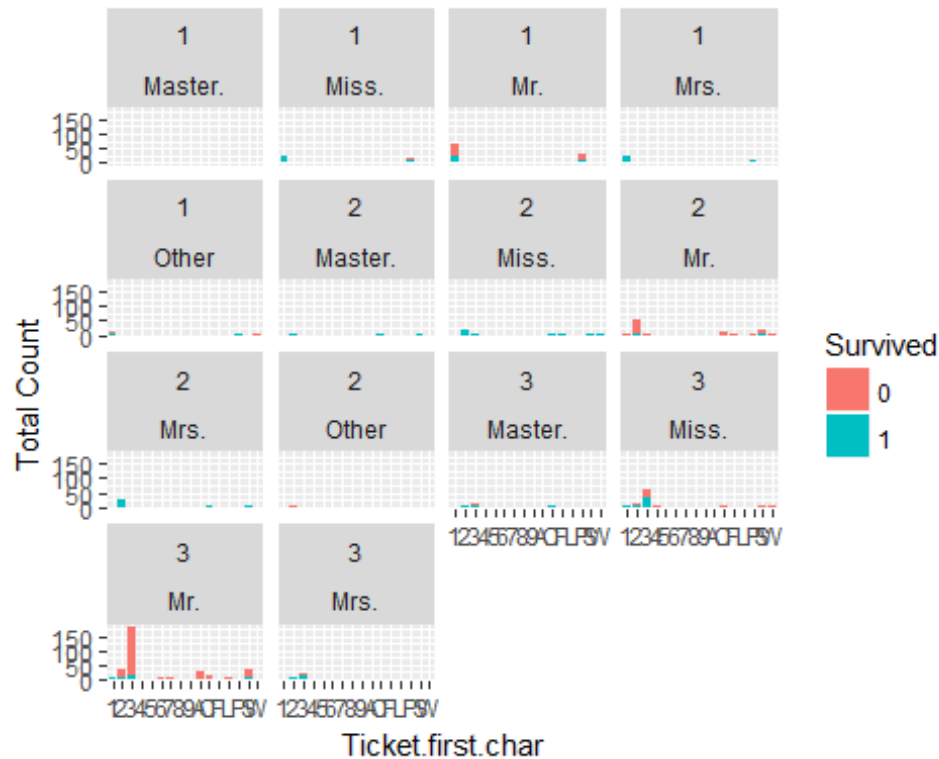
```
ggplot(data.combined[1:891,], aes(x = ticket.first.char, fill = Survived)) +
  geom_bar(width = 0.8) +
  ylab("Total Count") +
  xlab("Ticket.first.char") +
  labs(fill = "Survived")
```



```
# Using Pclass
ggplot(data.combined[1:891,], aes(x = ticket.first.char, fill = Survived)) +
  geom_bar(width = 0.8) +
  facet_wrap(~Pclass) +
  ylab("Total Count") +
  xlab("Ticket.first.char") +
  labs(fill = "Survived")
```



```
# Now using Pclass and Title both
ggplot(data.combined[1:891,], aes(x = ticket.first.char, fill = Survived))
+
  geom_bar(width = 0.8) +
  facet_wrap(~Pclass + title) +
  ylab("Total Count") +
  xlab("Ticket.first.char") +
  labs(fill = "Survived")
```



```
# Next up is Fare
summary(data.combined$Fare)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's 
##    0.000   7.896  14.454   33.295  31.275  512.329         1 

str(data.combined$Fare)    #numeric variable


##  num [1:1309] 7.25 71.28 7.92 53.1 8.05 ...

length(unique(data.combined$Fare))

## [1] 282


# We can relate the fair with Pclass
ggplot(data.combined, aes(x = Fare)) +
  geom_histogram(binwidth = 5) +
  xlab("Fare") +
  ylab("Total Count") +
  ylim(0,200)

## Warning: Removed 1 rows containing non-finite values (stat_bin).
## Warning: Removed 1 rows containing missing values (geom_bar).
```

```
# Let's see if it has some predictive power or not
ggplot(data.combined[1:891,], aes(x = Fare, fill = Survived)) +
  geom_histogram(binwidth = 20) +
  facet_wrap(~Pclass + title) +
  xlab("Fare") +
  ylab("Total Count") +
  ggtitle("Pclass + Title") +
  labs(fill = "Survived") +
  ylim(0,20)

## Warning: Removed 15 rows containing missing values (geom_bar).
```



Let's do something with cabin variable now

```
str(data.combined$Cabin)
```

```
## Factor w/ 187 levels "", "A10", "A14", ...: 1 83 1 57 1 1 131 1 1 1 ...
```

Clearly it's not a factor

```
data.combined$Cabin = as.character(data.combined$Cabin)
```

```
data.combined$Cabin[1:100]
```

```
## [1] "" "C85" "" "C123" ""
## [6] "" "E46" "" "" ""
## [11] "G6" "C103" "" "" ""
## [16] "" "" "" "" ""
## [21] "" "D56" "" "A6" ""
## [26] "" "" "C23 C25 C27" "" ""
## [31] "" "B78" "" "" ""
## [36] "" "" "" "" ""
## [41] "" "" "" "" ""
## [46] "" "" "" "" ""
## [51] "" "" "D33" "" "B30"
## [56] "C52" "" "" "" ""
## [61] "" "B28" "C83" "" ""
## [66] "" "F33" "" "" ""
## [71] "" "" "" "" ""
## [76] "F G73" "" "" "" ""
## [81] "" "" "" "" ""
## [86] "" "" "" "C23 C25 C27" ""
```

```
## [91] "" "" "E31" "" ""
## [96] "" "A5" "D10 D12" "" ""

# Replace the missing cabins with U
data.combined[which(data.combined$Cabin == ""), "Cabin"] = "U"
data.combined$Cabin[1:100]

## [1] "U" "C85" "U" "C123" "U"
## [6] "U" "E46" "U" "U" "U"
## [11] "G6" "C103" "U" "U" "U"
## [16] "U" "U" "U" "U" "U"
## [21] "U" "D56" "U" "A6" "U"
## [26] "U" "U" "C23 C25 C27" "U" "U"
## [31] "U" "B78" "U" "U" "U"
## [36] "U" "U" "U" "U" "U"
## [41] "U" "U" "U" "U" "U"
## [46] "U" "U" "U" "U" "U"
## [51] "U" "U" "D33" "U" "B30"
## [56] "C52" "U" "U" "U" "U"
## [61] "U" "B28" "C83" "U" "U"
## [66] "U" "F33" "U" "U" "U"
## [71] "U" "U" "U" "U" "U"
## [76] "F G73" "U" "U" "U" "U"
## [81] "U" "U" "U" "U" "U"
## [86] "U" "U" "U" "C23 C25 C27" "U"
## [91] "U" "U" "E31" "U" "U"
## [96] "U" "A5" "D10 D12" "U" "U"

# Take a Look at just first character or Letter
cabin.first.char = as.factor(substr(data.combined$Cabin,1,1))
str(cabin.first.char)

## Factor w/ 9 levels "A","B","C","D",...: 9 3 9 3 9 9 5 9 9 9 ...


levels(cabin.first.char)

## [1] "A" "B" "C" "D" "E" "F" "G" "T" "U"


# Adding it to combined data set and then we go on to plot it to see if
# there's any predictive power in it or not
data.combined$cabin.first.char = cabin.first.char

par(mfrow = c(3,3))

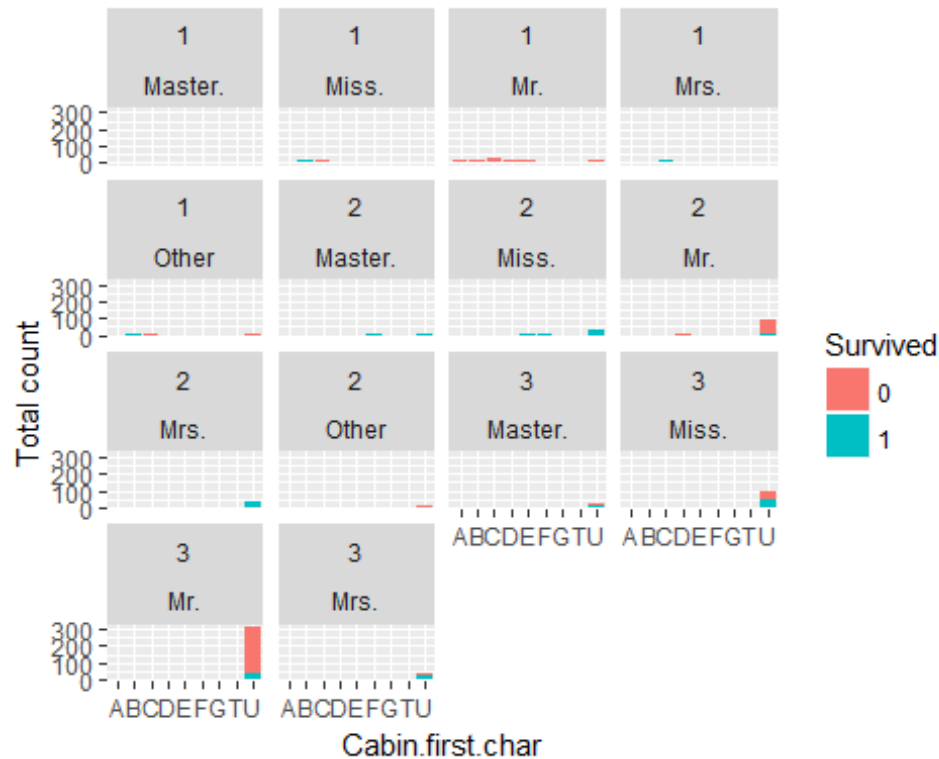
# Data Visualization again!
ggplot(data.combined[1:891,], aes(x = cabin.first.char, fill = Survived)) +
  geom_bar() +
  xlab("Cabin.first.char") +
  ylab("Total count")
```



```
#Let's drill in a bit more  
ggplot(data.combined[1:891,], aes(x = cabin.first.char, fill = Survived)) +  
  geom_bar() +  
  facet_wrap(~Pclass) +  
  xlab("Cabin.first.char") +  
  ylab("Total count")
```



```
#Pclass + title
ggplot(data.combined[1:891,], aes(x = cabin.first.char, fill = Survived)) +
  geom_bar() +
  facet_wrap(~Pclass + title) +
  xlab("Cabin.first.char") +
  ylab("Total count")
```




#What about folks with multiple cabins?

```
data.combined$cabin.multiple =  
as.factor(ifelse(str_detect(data.combined$Cabin, " "), "Y", "N"))
```

#Onto the ggplot thing, ofcourse.

```
ggplot(data.combined[1:891,], aes(x = cabin.multiple, fill = Survived)) +  
  geom_bar() +  
  facet_wrap(~Pclass + title)+  
xlab("Cabin.first.char") +  
ylab("Total count")
```



#Not particularly interesting. We shall come to it later on, maybe.

#Let's have a look at the last variable embarked

```
str(data.combined$Embarked)
```

```
## Factor w/ 4 levels "", "C", "Q", "S": 4 2 4 4 4 3 4 4 4 2 ...
```

```
summary(data.combined$Embarked)
```

```
##      C      Q      S
```

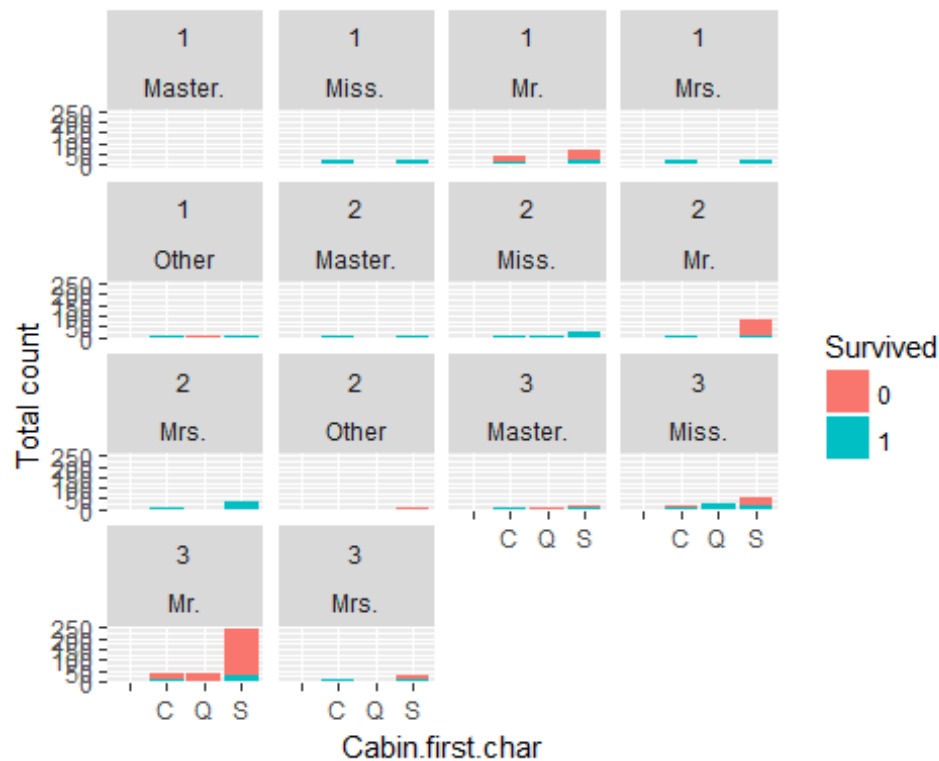
```
## 2 270 123 914
```

```
levels(data.combined$Embarked)
```

```
## [1] "" "C" "Q" "S"
```

#Some plotting again

```
ggplot(data.combined[1:891,], aes(x = Embarked, fill = Survived)) +
  geom_bar() +
  facet_wrap(~Pclass + title) +
  xlab("Cabin.first.char") +
  ylab("Total count")
```



Exploratory Data Analysis

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.4.2
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
```

```
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##      margin
```

#Let's train our first random forest model using just two predictor variables Pclass and title

```
rf.train.1 = data.combined[1:891, c("Pclass", "title")]
```

```
rf.label = as.factor(train$Survived)
```

```
set.seed(1234)
```

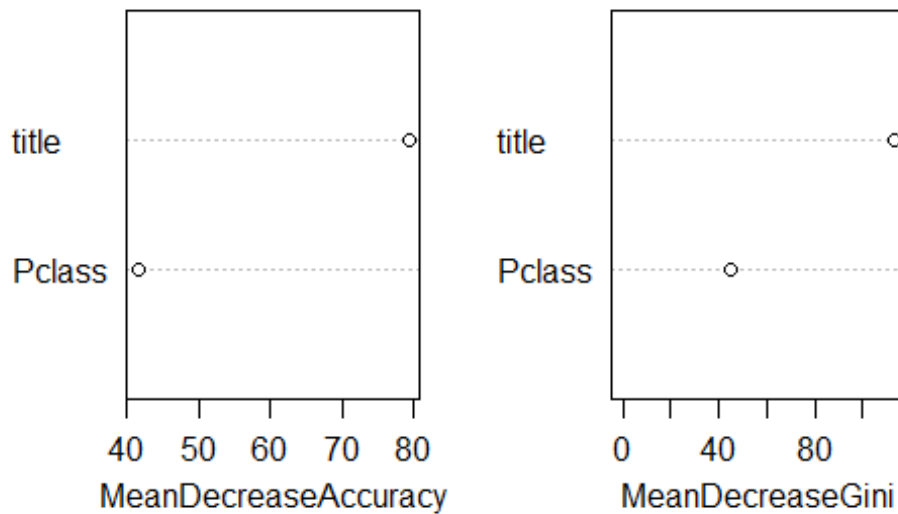
```
rf.1 = randomForest(x = rf.train.1, y = rf.label, importance = T, ntree = 1000)
```

```
rf.1
```



```
##
## Call:
## randomForest(x = rf.train.1, y = rf.label, ntree = 1000, importance = T)
##           Type of random forest: classification
##           Number of trees: 1000
## No. of variables tried at each split: 1
##
##           OOB estimate of  error rate: 20.76%
## Confusion matrix:
##      0   1 class.error
## 0 538  11  0.02003643
## 1 174 168  0.50877193
varImpPlot(rf.1)
```

rf.1



#Train a random forest model using Pclass, title and sibsp

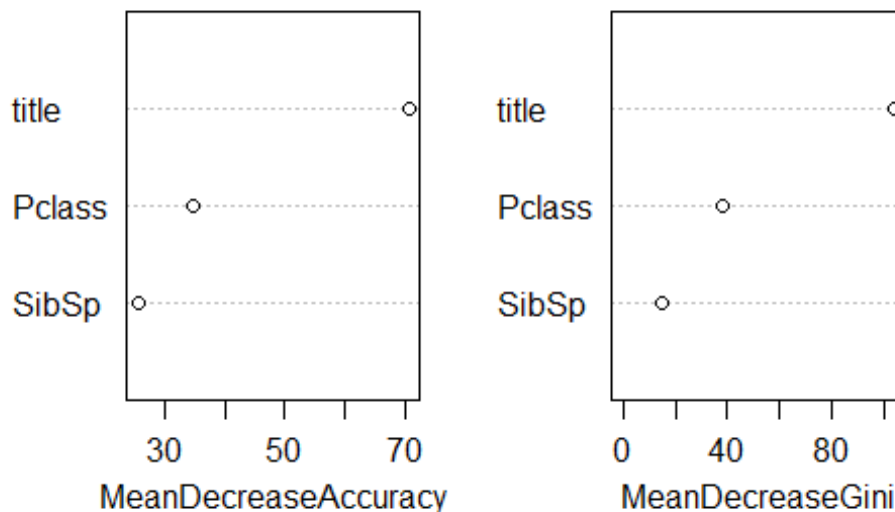
```
rf.train.2 = data.combined[1:891, c("Pclass", "title", "SibSp")]

set.seed(1234)
rf.2 = randomForest(x = rf.train.2, y = rf.label, importance = T, ntree = 1000)
rf.2

##
## Call:
## randomForest(x = rf.train.2, y = rf.label, ntree = 1000, importance = T)
```

```
##           Type of random forest: classification
##           Number of trees: 1000
## No. of variables tried at each split: 1
##
##           OOB estimate of  error rate: 19.75%
## Confusion matrix:
##      0   1 class.error
## 0 487  62   0.1129326
## 1 114 228   0.3333333
varImpPlot(rf.2)
```

rf.2



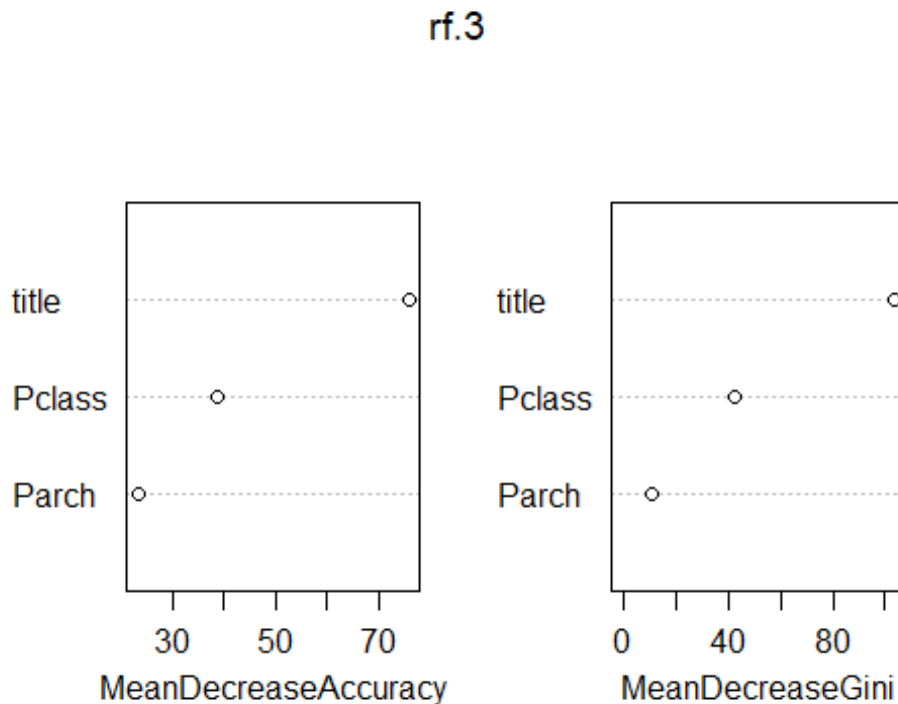
```
#Train a random forest model using Pclass, title and Parch

rf.train.3 = data.combined[1:891, c("Pclass", "title", "Parch")]

set.seed(1234)
rf.3 = randomForest(x = rf.train.3, y = rf.label, importance = T, ntree = 1000)
rf.3

##
## Call:
## randomForest(x = rf.train.3, y = rf.label, ntree = 1000, importance = T)
##           Type of random forest: classification
##           Number of trees: 1000
## No. of variables tried at each split: 1
```

```
##
##          OOB estimate of  error rate: 19.98%
## Confusion matrix:
##      0   1 class.error
## 0 495   54  0.09836066
## 1 124  218  0.36257310
varImpPlot(rf.3)
```



```
#Train a random forest model using Pclass, title, SibSp and Parch

rf.train.4 = data.combined[1:891, c("Pclass", "title", "Parch", "SibSp")]

set.seed(1234)
rf.4 = randomForest(x = rf.train.4, y = rf.label, importance = T, ntree = 1000)
rf.4

##
## Call:
## randomForest(x = rf.train.4, y = rf.label, ntree = 1000, importance = T)
##              Type of random forest: classification
##              Number of trees: 1000
## No. of variables tried at each split: 2
##
##          OOB estimate of  error rate: 18.63%
## Confusion matrix:
```

```
##      0    1 class.error
## 0 488  61   0.1111111
## 1 105 237   0.3070175
```

```
varImpPlot(rf.4)
```

rf.4



#Train a random forest model using Pclass, title and familysize

```
rf.train.5 = data.combined[1:891, c("Pclass", "title", "familysize")]
```

```
set.seed(1234)
```

```
rf.5 = randomForest(x = rf.train.5, y = rf.label, importance = T, ntree = 1000)
```

```
rf.5
```

```
##
```

```
## Call:
```

```
## randomForest(x = rf.train.5, y = rf.label, ntree = 1000, importance = T)
```

```
##           Type of random forest: classification
```

```
##           Number of trees: 1000
```

```
## No. of variables tried at each split: 1
```

```
##
```

```
##           OOB estimate of  error rate: 18.41%
```

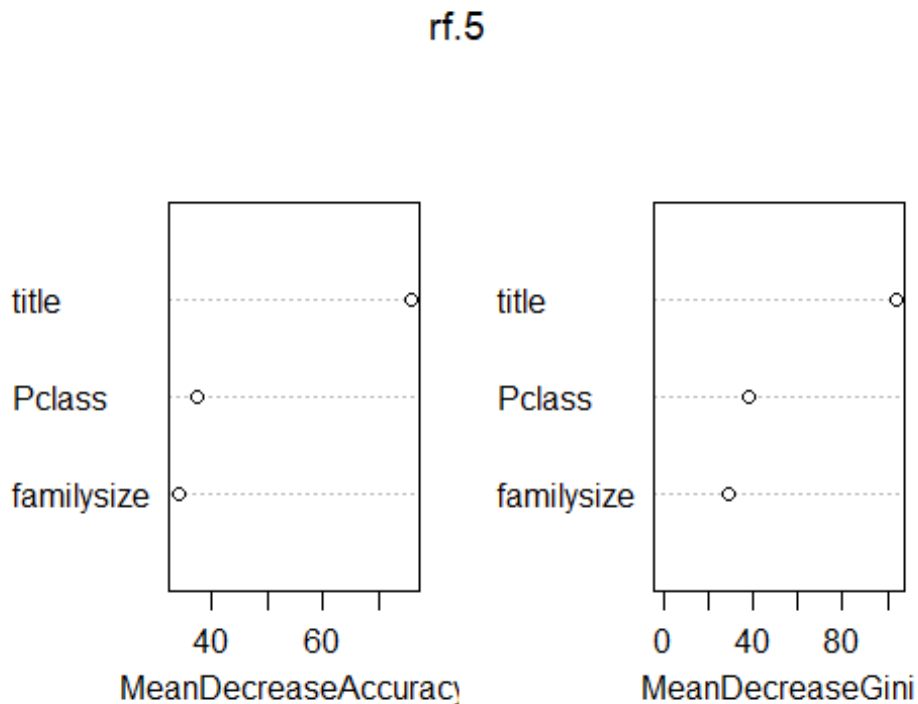
```
## Confusion matrix:
```

```
##      0    1 class.error
```

```
## 0 485  64   0.1165756
```

```
## 1 100 242   0.2923977
```

```
varImpPlot(rf.5)
```



```
#Train a random forest model using Pclass, title, familysize and Parch
```

```
rf.train.6 = data.combined[1:891, c("Pclass", "title", "Parch",  
"familysize")]
```

```
set.seed(1234)
```

```
rf.6 = randomForest(x = rf.train.6, y = rf.label, importance = T, ntree =  
1000)  
rf.6
```

```
##
```

```
## Call:
```

```
## randomForest(x = rf.train.6, y = rf.label, ntree = 1000, importance = T)
```

```
##           Type of random forest: classification
```

```
##           Number of trees: 1000
```

```
## No. of variables tried at each split: 2
```

```
##
```

```
##           OOB estimate of  error rate: 18.97%
```

```
## Confusion matrix:
```


```
##      0   1 class.error
```

```
## 0 486  63  0.1147541
```

```
## 1 106 236  0.3099415
```

```
varImpPlot(rf.6)
```

rf.6



```
#Train a random forest model using Pclass, title, SibSp and Familysize
```

```
rf.train.7 = data.combined[1:891, c("Pclass", "title", "SibSp",  
"familysize")]
```

```
set.seed(1234)
```

```
rf.7 = randomForest(x = rf.train.7, y = rf.label, importance = T, ntree =  
1000)
```

```
rf.7
```

```
##
```

```
## Call:
```

```
## randomForest(x = rf.train.7, y = rf.label, ntree = 1000, importance = T)
```

```
## Type of random forest: classification
```

```
## Number of trees: 1000
```

```
## No. of variables tried at each split: 2
```

```
##
```

```
## OOB estimate of error rate: 18.74%
```

```
## Confusion matrix:
```


```
## 0 1 class.error
```

```
## 0 486 63 0.1147541
```

```
## 1 104 238 0.3040936
```

```
varImpPlot(rf.7)
```

rf.7



#As it is very evident from the above running models, the model with just title, pclass and familysize gives the highest accuracy or Lowest OOB rate

Cross Validation

#Let's try to submit these predictions to Kaggle first and see how we're doing

```
test.submit.df = data.combined[892:1309, c("Pclass", "familysize", "title")]
```

#This is how you should predict

```
rf.5.preds = predict(rf.5, test.submit.df)
table(rf.5.preds)
```

```
## rf.5.preds
```

```
##  0  1
```

```
## 258 160
```

#Write out a CSV file for the submission to Kaggle

```
submit.df = data.frame(PassengerId = 892:1309, Survived = rf.5.preds)
```

```
write.csv(submit.df, file = "RF1.csv", row.names = F)
```

#Now, as we can see from the Kaggle, our score turns out be 0.79426 but the OOB estimates predicted it to be 0.8159

#Let's dig deep into the concept of cross-validation

```
library(caret)
```

```

## Warning: package 'caret' was built under R version 3.4.4
## Loading required package: lattice
library(doSNOW)
## Warning: package 'doSNOW' was built under R version 3.4.4
## Loading required package: foreach
## Warning: package 'foreach' was built under R version 3.4.3
## Loading required package: iterators
## Warning: package 'iterators' was built under R version 3.4.3
## Loading required package: snow
## Warning: package 'snow' was built under R version 3.4.4
#We're gonna be doing something called stratified cross-validation.
set.seed(2348)
cv.10.folds = createMultiFolds(y=rf.label, k=10, times = 10)

#Check stratification
table(rf.label)

## rf.label
##    0    1
## 549 342

342/549

## [1] 0.6229508

#Check for any fold now
table(rf.label[cv.10.folds[[33]]])

##
##    0    1
## 494 308

307/494

## [1] 0.6214575

#For stratification, the main property is that the ratio of folks who
perished to the folks who survived should be same in the each folds and the
y(rf.label)

#Now, let's setup traincontrol object per above
ctrl.1 = trainControl(method = "repeatedcv", number = 10, repeats = 10, index
= cv.10.folds)

```



```
#Set up doSNOW package for multi-core training. This is helpful because we're gonna be using a lot of trees
cl = makeCluster(6, type = "SOCK")
registerDoSNOW(cl)
```

```
#Set seed for reproducibility and train
set.seed(34324)
rf.5.cv.1 = train(x=rf.train.5, y=rf.label, method="rf", tuneLength=3,
ntree=1000, trControl= ctrl.1)
```

```
## note: only 2 unique complexity parameters in default grid. Truncating the grid to 2 .
```

```
#Shut down cluster
stopCluster(cl)
```

```
#Check out results
rf.5.cv.1
```

```
## Random Forest
##
## 891 samples
## 3 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 10 times)
## Summary of sample sizes: 801, 802, 802, 803, 802, 801, ...
## Resampling results across tuning parameters:
##
##  mtry  Accuracy   Kappa
##  2      0.8128058 0.594528
##  3      0.8093388 0.585973
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

```
#Let's also try 5-fold to see if there is improvment in accuracy
```

```
set.seed(5983)
cv.5.folds = createMultiFolds(y=rf.label, k=5, times = 10)
```

```
ctrl.2 = trainControl(method = "repeatedcv", number = 5, repeats = 10, index
= cv.5.folds)
```

```
#Set up doSNOW package for multi-core training. This is helpful because we're gonna be using a lot of trees
cl = makeCluster(6, type = "SOCK")
registerDoSNOW(cl)
```

```
#Set seed for reproducibility and train
```

```

set.seed(89472)
rf.5.cv.2 = train(x=rf.train.5, y=rf.label, method="rf", tuneLength=3,
ntree=1000, trControl= ctrl.2)

## note: only 2 unique complexity parameters in default grid. Truncating the
grid to 2 .

#Shut down cluster
stopCluster(cl)

#Check out results
rf.5.cv.2

## Random Forest
##
## 891 samples
## 3 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold, repeated 10 times)
## Summary of sample sizes: 713, 713, 713, 713, 712, 713, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.8133608 0.5974520
## 3 0.8093159 0.5881247
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.

#3-fold
set.seed(5986)
cv.3.folds = createMultiFolds(y=rf.label, k=3, times = 10)

ctrl.3 = trainControl(method = "repeatedcv", number = 3, repeats = 10, index
= cv.3.folds)

#Set up doSNOW package for multi-core training. This is helpful because we're
gonna be using a lot of trees
cl = makeCluster(6, type = "SOCK")
registerDoSNOW(cl)

#Set seed for reproducibility and train
set.seed(89465)
rf.5.cv.3 = train(x=rf.train.5, y=rf.label, method="rf", tuneLength=3,
ntree=1000, trControl= ctrl.3)

## note: only 2 unique complexity parameters in default grid. Truncating the
grid to 2 .

```

```

#Shut down cluster
stopCluster(cl)

#Check out results
rf.5.cv.3

## Random Forest
##
## 891 samples
## 3 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (3 fold, repeated 10 times)
## Summary of sample sizes: 594, 594, 594, 594, 594, 594, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.8143659 0.5956886
## 3 0.8104377 0.5857806
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.

#Let's see where we might have gone wrong. Let's build a single decision tree
to check what exactly is happening on the inside
#Random forests are ofcourse way better than decision trees but when it comes
to easily understand the whole picture, deciion trees
#are way better than random forests

library(rpart)
#install.packages("rpart.plot")
library(rpart.plot)

## Warning: package 'rpart.plot' was built under R version 3.4.4

#Create utility function
rpart.cv = function(seed, training, labels, ctrl) {
  cl = makeCluster(6, type = "SOCK")
  registerDoSNOW(cl)

  set.seed(seed)
  #Leverage formula interface for training
  rpart.cv = train(x=training, y=labels, method="rpart", tuneLength=30,
trControl = ctrl)

  #shutdown cluster
  stopCluster(cl)

  return(rpart.cv)

```

```
}
```

```
#Grab features
```

```
features = c("Pclass", "title", "familysize")  
rpart.train.1 = data.combined[1:891, features]
```

```
#Run CV and check out results
```

```
rpart.1.cv.1 = rpart.cv(94622, rpart.train.1, rf.label, ctrl.3)  
rpart.1.cv.1
```

```
## CART
```

```
##
```

```
## 891 samples
```

```
## 3 predictor
```

```
## 2 classes: '0', '1'
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (3 fold, repeated 10 times)
```

```
## Summary of sample sizes: 594, 594, 594, 594, 594, 594, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

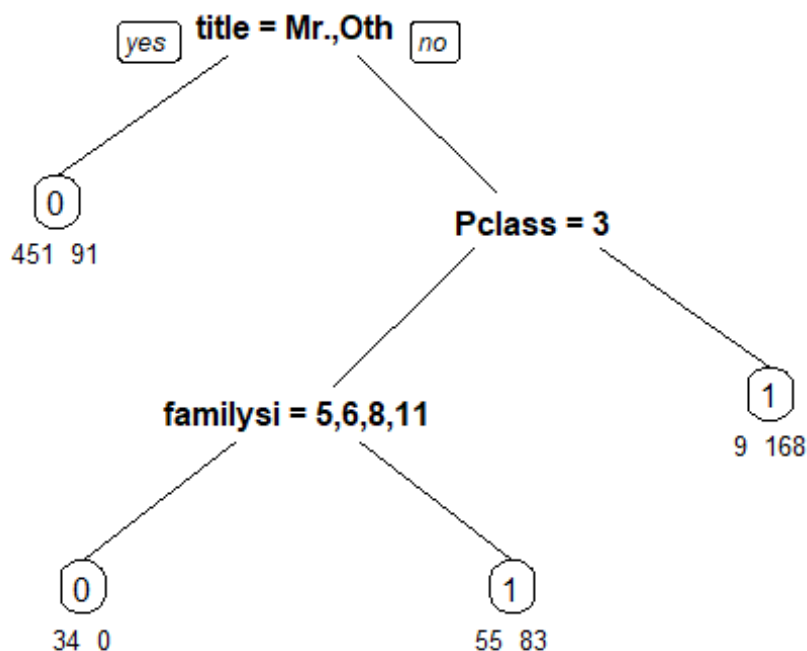
##	cp	Accuracy	Kappa
##	0.00000000	0.8116723	0.5917680
##	0.01542650	0.8205387	0.6145781
##	0.03085299	0.8205387	0.6145781
##	0.04627949	0.8129068	0.5999912
##	0.06170599	0.7888889	0.5527336
##	0.07713249	0.7888889	0.5527336
##	0.09255898	0.7867565	0.5511186
##	0.10798548	0.7848485	0.5477607
##	0.12341198	0.7842873	0.5469939
##	0.13883848	0.7842873	0.5469939
##	0.15426497	0.7842873	0.5469939
##	0.16969147	0.7842873	0.5469939
##	0.18511797	0.7842873	0.5469939
##	0.20054446	0.7842873	0.5469939
##	0.21597096	0.7842873	0.5469939
##	0.23139746	0.7842873	0.5469939
##	0.24682396	0.7842873	0.5469939
##	0.26225045	0.7842873	0.5469939
##	0.27767695	0.7842873	0.5469939
##	0.29310345	0.7842873	0.5469939
##	0.30852995	0.7842873	0.5469939
##	0.32395644	0.7842873	0.5469939
##	0.33938294	0.7842873	0.5469939
##	0.35480944	0.7842873	0.5469939
##	0.37023593	0.7842873	0.5469939
##	0.38566243	0.7842873	0.5469939
##	0.40108893	0.7777778	0.5267317
##	0.41651543	0.7582492	0.4661039

```
## 0.43194192 0.7202020 0.3472667
## 0.44736842 0.6868687 0.2382834
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.03085299.
```

#Plot

```
prp(rpart.1.cv.1$finalModel, type = 0, extra = 1, under = T)
```

```
## Warning: Bad 'data' field in model 'call' field.
##      To make this warning go away:
##      Call prp with roundint=FALSE,
##      or rebuild the rpart model with model=TRUE.
```



```
#The plot brings out some interesting lines of investigation. Namely:
# 1 - Titles of Mr. and Others are predicted to perish at an overall
accuracy of 83.2%
# 2 - Titles of Master, Miss and Mrs. in 1st and 2nd class are predicted
to survive at
#      an overall accuracy of 94.9%
# 3 - Titles of Master, Miss & Mrs in class 3 and having a family size of
5,6,8,11
#      are predicted to perish at an overall accuracy of 100%
# 4 - Titles of Master, Miss & Mrs in class 3 and having a family size of
5,6,8,11
#      are predicted to survive at an overall accuracy of 59.6%
```

```

#Both rpart and ef confirm that title is important. Let's investigate further:
#Also, we're stressing more on the 1st point here that the title Mr and other just seem blunt.
#Let's move ahead and investigate it further
table(data.combined$title)

##
## Master.    Miss.      Mr.      Mrs.    Other
##         61      260      758      199      31

#Parse out last name and title
data.combined$Name[1:5]

## [1] Braund, Mr. Owen Harris
## [2] Cumings, Mrs. John Bradley (Florence Briggs Thayer)
## [3] Heikkinen, Miss. Laina
## [4] Futrelle, Mrs. Jacques Heath (Lily May Peel)
## [5] Allen, Mr. William Henry
## 1307 Levels: Abbing, Mr. Anthony ... Zakarian, Mr. Ortin

name.splits = str_split(data.combined$Name, ",")
name.splits[1]

## [[1]]
## [1] "Braund"          " Mr. Owen Harris"

last.names = sapply(name.splits, "[", 1)
last.names[1:10]

## [1] "Braund"    "Cumings"   "Heikkinen" "Futrelle"  "Allen"
## [6] "Moran"     "McCarthy"  "Palsson"   "Johnson"  "Nasser"

#Add last names to the data.combined in case we might find it useful later
data.combined$last.name = last.names

#Now for titles
name.splits = str_split(sapply(name.splits, "[", 2), " ")
titles = sapply(name.splits, "[", 2)
unique(titles)

## [1] "Mr."      "Mrs."     "Miss."    "Master."  "Don."
## [6] "Rev."     "Dr."      "Mme."     "Ms."      "Major."
## [11] "Lady."   "Sir."     "Mlle."    "Col."     "Capt."
## [16] "the"      "Jonkheer." "Dona."

#What's up with the title of "the"?
which(titles == "the")

## [1] 760

data.combined[760,]

```

```
##      PassengerId Survived Pclass
## 760          760          1      1
##
##                                     Name      Sex Age
## 760  Rothes, the Countess. of (Lucy Noel Martha Dyer-Edwards) female  33
##      SibSp Parch Ticket Fare Cabin Embarked title familysize
## 760      0      0 110152 86.5   B77          S Other          1
##      ticket.first.char cabin.first.char cabin.multiple last.name
## 760                  1                  B                  N    Rothes
```

#Re-map titles to be more exact

```
titles[titles %in% c("the", "Dona.")] = "Lady."
titles[titles %in% c("Ms.", "Mlle.")] = "Miss."
titles[titles == "Mme."] = "Mrs."
titles[titles %in% c("Jonkheer", "Don.")] = "Sir."
titles[titles %in% c("Col.", "Capt.", "Major.")] = "Officer"
table(titles)
```


```
## titles
##      Dr. Jonkheer.      Lady.      Master.      Miss.      Mr.      Mrs.
##      8          1          3          61         264         757         198
## Officer      Rev.      Sir.
##      7          8          2
```

#Now add this to the dataframe again

```
data.combined$new.title = as.factor(titles)
```

#Let's again use this for data visualization

```
ggplot(data.combined[1:891,], aes(x=new.title, fill=Survived)) +
  geom_bar(width = 0.5) +
  facet_wrap(~Pclass) +
  xlab("New title") +
  ylab("Counts") +
  labs(fill = "Survived")
```




#Collapse titles based on visual analysis

```
indexes = which(data.combined$new.title == "Lady.")
data.combined$new.title[indexes] = "Mrs."
```

```
indexes = which(data.combined$new.title == "Dr." |
                 data.combined$new.title == "Rev." |
                 data.combined$new.title == "Sir." |
                 data.combined$new.title == "Officer")
data.combined$new.title[indexes] = "Mr."
```

#Visualize again

```
ggplot(data.combined[1:891,], aes(x=new.title, fill=Survived)) +
  geom_bar(width = 0.5) +
  facet_wrap(~Pclass) +
  xlab("New title") +
  ylab("Counts") +
  labs(fill = "Survived")
```

#Grab features

```
features = c("Pclass", "new.title", "familysize")
rpart.train.2 = data.combined[1:891, features]
```

#Run CV and check out results

```
rpart.2.cv.1 = rpart.cv(94622, rpart.train.2, rf.label, ctrl.3)
rpart.2.cv.1
```

CART

##

891 samples

3 predictor

2 classes: '0', '1'

##

No pre-processing

Resampling: Cross-Validated (3 fold, repeated 10 times)

Summary of sample sizes: 594, 594, 594, 594, 594, 594, ...

Resampling results across tuning parameters:

##

##	cp	Accuracy	Kappa
##	0.00000000	0.8191919	0.6070484
##	0.01582980	0.8287318	0.6319625
##	0.03165961	0.8287318	0.6319625
##	0.04748941	0.8210999	0.6173617
##	0.06331922	0.7959596	0.5677607
##	0.07914902	0.7973064	0.5721536
##	0.09497883	0.7938272	0.5660315

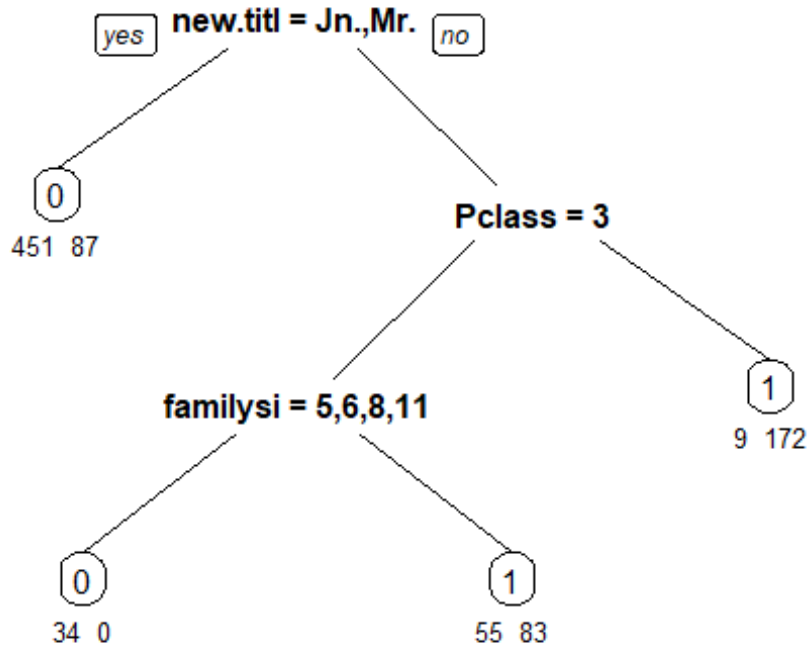
```

## 0.11080863 0.7929293 0.5643804
## 0.12663844 0.7923681 0.5636067
## 0.14246824 0.7923681 0.5636067
## 0.15829804 0.7923681 0.5636067
## 0.17412785 0.7923681 0.5636067
## 0.18995765 0.7923681 0.5636067
## 0.20578746 0.7923681 0.5636067
## 0.22161726 0.7923681 0.5636067
## 0.23744707 0.7923681 0.5636067
## 0.25327687 0.7923681 0.5636067
## 0.26910667 0.7923681 0.5636067
## 0.28493648 0.7923681 0.5636067
## 0.30076628 0.7923681 0.5636067
## 0.31659609 0.7923681 0.5636067
## 0.33242589 0.7923681 0.5636067
## 0.34825570 0.7923681 0.5636067
## 0.36408550 0.7923681 0.5636067
## 0.37991531 0.7923681 0.5636067
## 0.39574511 0.7923681 0.5636067
## 0.41157491 0.7850730 0.5419190
## 0.42740472 0.7645342 0.4793351
## 0.44323452 0.7193042 0.3386920
## 0.45906433 0.7012346 0.2816186
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.03165961.

#Plot
prp(rpart.2.cv.1$finalModel, type = 0, extra = 1, under = T)

## Warning: Bad 'data' field in model 'call' field.
##       To make this warning go away:
##       Call prp with roundint=FALSE,
##       or rebuild the rpart model with model=TRUE.

```



#Dive in on 1st class Mr.

```
indexes.first.mr = which(data.combined$new.title == "Mr." &
data.combined$Pclass == "1")
first.mr.df = data.combined[indexes.first.mr,]
summary(first.mr.df)
```

```
## PassengerId      Survived  Pclass
## Min.   :  7.0      0   :76    1:174
## 1st Qu.: 371.8      1   :43    2:  0
## Median : 647.0     None:55    3:  0
## Mean   : 656.8
## 3rd Qu.: 966.5
## Max.   :1299.0
##
##                                     Name      Sex      Age
## Anderson, Mr. Harry                :  1  female:  1  Min.   :17.00
## Andrews, Mr. Thomas Jr              :  1   male :173  1st Qu.:31.00
## Artagaveytia, Mr. Ramon              :  1                                     Median :42.00
## Barkworth, Mr. Algernon Henry Wilson:  1                                     Mean    :42.27
## Baumann, Mr. John D                  :  1                                     3rd Qu.:50.75
## Baxter, Mr. Quigg Edmond              :  1                                     Max.    :80.00
## (Other)                             :168                                     NA's    :28
## SibSp      Parch      Ticket      Fare
## 0:121  0      :145  Length:174    Min.   :  0.00
## 1: 50  1      : 21  Class :character 1st Qu.: 27.72
## 2:  2  2      :  6  Mode  :character Median : 46.30
## 3:  1  3      :  1                                     Mean   : 67.73
```

```
## 4: 0 4 : 1 3rd Qu.: 78.46
## 5: 0 5 : 0 Max. :512.33
## 8: 0 (Other): 0
## Cabin Embarked title familysize ticket.first.char
## Length:174 : 0 Master.: 0 1 :108 1 :109
## Class :character C: 68 Miss. : 0 2 : 45 P : 45
## Mode :character Q: 1 Mr. :160 3 : 16 3 : 8
## S:105 Mrs. : 0 4 : 2 6 : 4
## Other : 14 6 : 2 2 : 3
## 5 : 1 F : 2
## (Other): 0 (Other): 3
## cabin.first.char cabin.multiple last.name new.title
## C :47 N:161 Length:174 Mr. :174
## U :43 Y: 13 Class :character Dr. : 0
## B :27 Mode :character Jonkheer.: 0
## D :20 Lady. : 0
## E :19 Master. : 0
## A :17 Miss. : 0
## (Other): 1 (Other) : 0
```

#1 female?

```
first.mr.df[first.mr.df$Sex == "female",]
```

```
## PassengerId Survived Pclass Name Sex Age
## 797 797 1 1 Leader, Dr. Alice (Farnham) female 49
## SibSp Parch Ticket Fare Cabin Embarked title familysize
## 797 0 0 17465 25.9292 D17 S Other 1
## ticket.first.char cabin.first.char cabin.multiple last.name new.title
## 797 1 D N Leader Mr.
```

#Here, we can see that she is Dr. and has been classified as a "Mr."

#Let's update new.title feature

```
indexes = which(data.combined$Sex == "female" & data.combined$new.title ==
"Mr.")
data.combined$new.title[indexes] = "Mrs."
```

#Any other gender slip-ups?

```
length(which(data.combined$Sex == "female" & (data.combined$new.title ==
"Master." | data.combined$new.title == "Mr.")))
```

```
## [1] 0
```

#Refresh dataframe

```
indexes.first.mr = which(data.combined$new.title == "Mr." &
data.combined$Pclass == "1")
first.mr.df = data.combined[indexes.first.mr,]
```

#Let's look at surviving 1st class "Mr."

```
summary(first.mr.df[first.mr.df$Survived == "1",])
```

```
## PassengerId Survived Pclass
## Min. : 24.0 0 : 0 1:42
## 1st Qu.:435.2 1 :42 2: 0
## Median :594.0 None: 0 3: 0
## Mean :528.5
## 3rd Qu.:681.5
## Max. :890.0
##
## Name Sex
## Anderson, Mr. Harry : 1 female: 0
## Barkworth, Mr. Algernon Henry Wilson : 1 male :42
## Beckwith, Mr. Richard Leonard : 1
## Behr, Mr. Karl Howell : 1
## Bishop, Mr. Dickinson H : 1
## Bjornstrom-Steffansson, Mr. Mauritz Hakan: 1
## (Other) :36
## Age SibSp Parch Ticket Fare
## Min. :17.00 0:28 0 :36 Length:42 Min. : 26.29
## 1st Qu.:28.00 1:13 1 : 4 Class :character 1st Qu.: 27.34
## Median :36.00 2: 1 2 : 2 Mode :character Median : 35.50
## Mean :38.76 3: 0 3 : 0 Mean : 71.55
## 3rd Qu.:48.00 4: 0 4 : 0 3rd Qu.: 73.39
## Max. :80.00 5: 0 5 : 0 Max. :512.33
## NA's :5 8: 0 (Other): 0
## Cabin Embarked title familysize ticket.first.char
## Length:42 : 0 Master.: 0 1 :25 1 :30
## Class :character C:17 Miss. : 0 2 :12 P :11
## Mode :character Q: 0 Mr. :38 3 : 4 2 : 1
## S:25 Mrs. : 0 4 : 1 3 : 0
## Other : 4 5 : 0 4 : 0
## 6 : 0 5 : 0
## (Other): 0 (Other): 0
## cabin.first.char cabin.multiple last.name new.title
## C :10 N:39 Length:42 Mr. :42
## E : 8 Y: 3 Class :character Dr. : 0
## B : 7 Mode :character Jonkheer.: 0
## D : 6 Lady. : 0
## U : 6 Master. : 0
## A : 5 Miss. : 0
## (Other): 0 (Other) : 0
```

```
View(first.mr.df[first.mr.df$Survived == "1",])
```

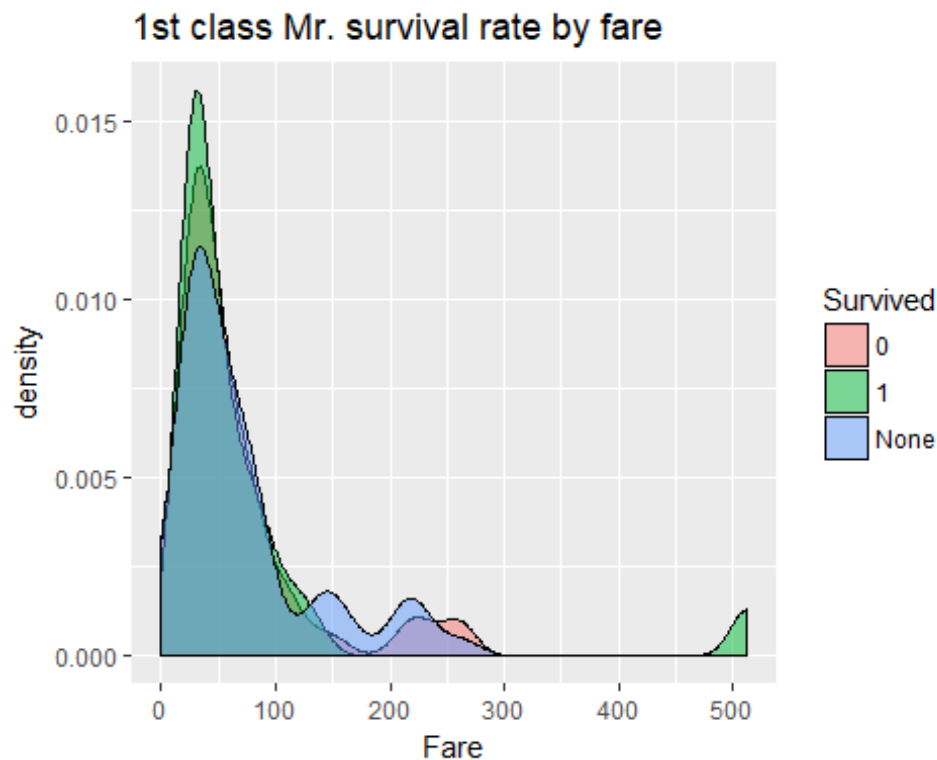
#Take a Look at some of the high fares

```
indexes = which(data.combined$Ticket == "PC 17755" |
                data.combined$Ticket == "113760" |
                data.combined$Ticket == "PC 17611")
```

```
View(data.combined[indexes,])
```

#Visualize survival rates for 1st class "Mr." by fare

```
ggplot(first.mr.df, aes(x = Fare, fill = Survived)) +
  geom_density(alpha = 0.5) +
  ggtitle("1st class Mr. survival rate by fare")
```



```
#Engineer features based on all the passengers with the same ticket
ticket.party.size = rep(0, nrow(data.combined))
avg.fare = rep(0.0, nrow(data.combined) )
tickets = unique(data.combined$Ticket)

for(i in 1:length(tickets)) {
  current.ticket = tickets[i]
  party.indexes = which(data.combined$Ticket == current.ticket)
  current.avg.fare = data.combined[party.indexes[1], "Fare"] /
length(party.indexes)

  for(k in i:length(party.indexes)){
    ticket.party.size[party.indexes[k]] = length(party.indexes)
    avg.fare[party.indexes[k]] = current.avg.fare
  }
}

data.combined$ticket.party.size = ticket.party.size
data.combined$avg.fare = avg.fare
data.combined$`avg. fare` = NULL

#Refresh 1st class "Mr." dataframe
```

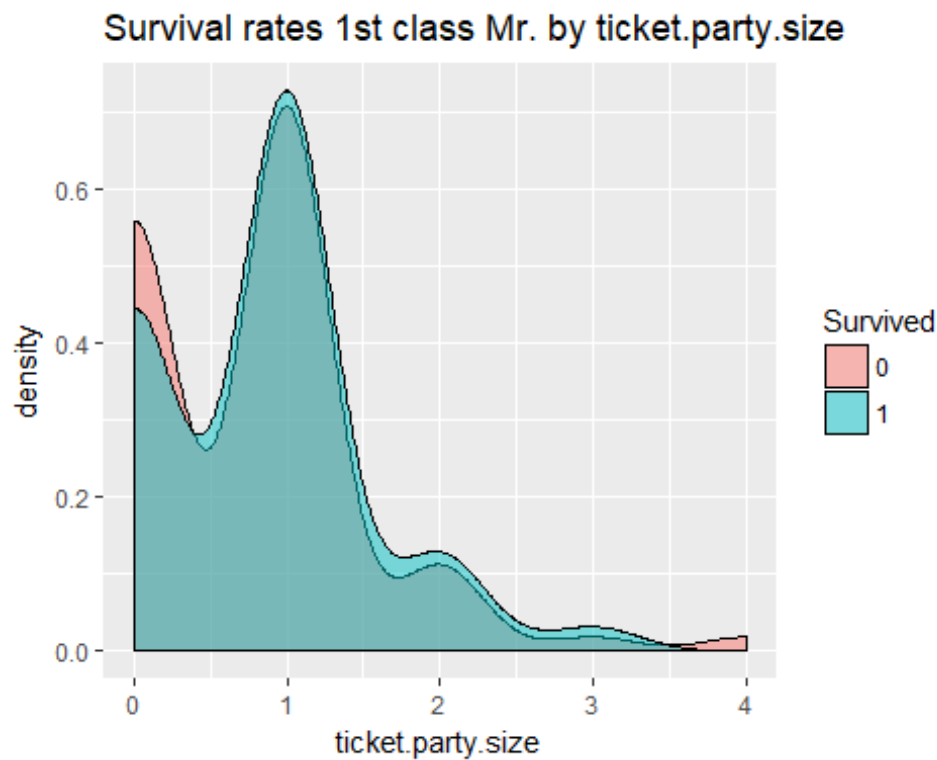
```
first.mr.df = data.combined[indexes.first.mr,]
summary(first.mr.df)
```

```
## PassengerId Survived Pclass
## Min. : 7 0 :76 1:173
## 1st Qu.: 371 1 :42 2: 0
## Median : 646 None:55 3: 0
## Mean : 656
## 3rd Qu.: 967
## Max. :1299
##
## Name Sex Age
## Anderson, Mr. Harry : 1 female: 0 Min. :17.00
## Andrews, Mr. Thomas Jr : 1 male :173 1st Qu.:31.00
## Artagaveytia, Mr. Ramon : 1 Median :42.00
## Barkworth, Mr. Algernon Henry Wilson: 1 Mean :42.22
## Baumann, Mr. John D : 1 3rd Qu.:51.00
## Baxter, Mr. Quigg Edmond : 1 Max. :80.00
## (Other) :167 NA's :28
## SibSp Parch Ticket Fare
## 0:120 0 :144 Length:173 Min. : 0.00
## 1: 50 1 : 21 Class :character 1st Qu.: 27.72
## 2: 2 2 : 6 Mode :character Median : 47.10
## 3: 1 3 : 1 Mean : 67.98
## 4: 0 4 : 1 3rd Qu.: 78.85
## 5: 0 5 : 0 Max. :512.33
## 8: 0 (Other): 0
## Cabin Embarked title familysize ticket.first.char
## Length:173 : 0 Master.: 0 1 :107 1 :108
## Class :character C: 68 Miss. : 0 2 : 45 P : 45
## Mode :character Q: 1 Mr. :160 3 : 16 3 : 8
## S:104 Mrs. : 0 4 : 2 6 : 4
## Other : 13 6 : 2 2 : 3
## 5 : 1 F : 2
## (Other): 0 (Other): 3
## cabin.first.char cabin.multiple last.name new.title
## C :47 N:160 Length:173 Mr. :173
## U :43 Y: 13 Class :character Dr. : 0
## B :27 Mode :character Jonkheer.: 0
## D :19 Lady. : 0
## E :19 Master. : 0
## A :17 Miss. : 0
## (Other): 1 (Other) : 0
## ticket.party.size avg.fare
## Min. :0.0000 Min. : 0.00
## 1st Qu.:0.0000 1st Qu.: 0.00
## Median :1.0000 Median :26.55
## Mean :0.9711 Mean :20.52
## 3rd Qu.:1.0000 3rd Qu.:30.50
```

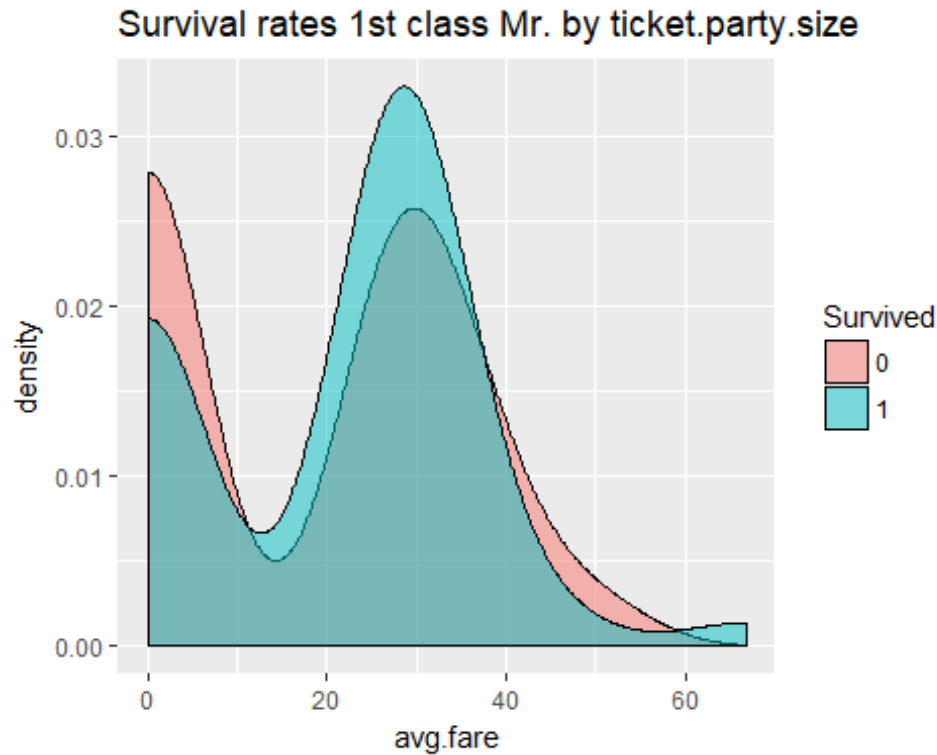
```
## Max. :6.0000 Max. :66.83
##
```

#Visualize new features

```
ggplot(first.mr.df[first.mr.df$Survived != "None",], aes(x =  
ticket.party.size, fill=Survived)) +  
  geom_density(alpha = 0.5) +  
  ggtitle("Survival rates 1st class Mr. by ticket.party.size")
```



```
ggplot(first.mr.df[first.mr.df$Survived != "None",], aes(x = avg.fare,  
fill=Survived)) +  
  geom_density(alpha = 0.5) +  
  ggtitle("Survival rates 1st class Mr. by ticket.party.size")
```

#Hypothesis - ticket.party.size is highly correlated with avg. fare

```
summary(data.combined$avg.fare)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##      0.000  0.000   7.775   9.513  10.500  128.082      1
```

#Let's figure out the NA value

```
which(is.na(data.combined$avg.fare))
```

```
## [1] 1044
```

```
data.combined[1044,]
```

```
##      PassengerId Survived Pclass      Name Sex Age SibSp Parch
## 1044         1044     None      3 Storey, Mr. Thomas male 60.5    0    0
##      Ticket Fare Cabin Embarked title familysize ticket.first.char
## 1044   3701   NA      U          S Mr.          1              3
##      cabin.first.char cabin.multiple last.name new.title ticket.party.size
## 1044                U              N   Storey      Mr.              1
##      avg.fare
## 1044      NA
```

#Get records for similar passengers and summarize avg. fares

```
indexes = with(data.combined, which(Pclass == "3" & title == "Mr." &
familysize == "1" & Ticket != "3701"))
similar.na.passengers = data.combined[indexes,]
summary(similar.na.passengers$avg.fare)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.000   7.250   7.798   7.359   8.050   10.171

#Use median since it is very close to mean and slightly higher than mean
data.combined[is.na(avg.fare), "avg.fare"] = 7.840

#Leverage caret's preProcess function to normalize data
prepoc.data.combined = data.combined[,c("ticket.party.size", "avg.fare")]
prePoc = preProcess(preproc.data.combined, method = c("center", "scale"))

postproc.data.combined = predict(prePoc, prepoc.data.combined)

#Let's check the correlation between avg.fare and ticket.party.size
cor(postproc.data.combined$ticket.party.size,
postproc.data.combined$avg.fare)

## [1] 0.4485116

#Correlation always results between -1 and 1 where -1 is negatively
correlated and +1 means completely correlated. 0 means no correlation at all
#Here, they are highly uncorrelated means we have two new potential features
that we could add

#How about just 1st class all up?
indexes = which(data.combined$Pclass == "1")
cor(postproc.data.combined$ticket.party.size[indexes],
postproc.data.combined$avg.fare[indexes])

## [1] 0.7088304

#Hypothesis refuted again

#Okay. Let's see if our feature engineering has made any difference or not
features = c("Pclass", "new.title", "familysize", "ticket.party.size",
"avg.fare")
rpart.train.3 = data.combined[1:891, features]

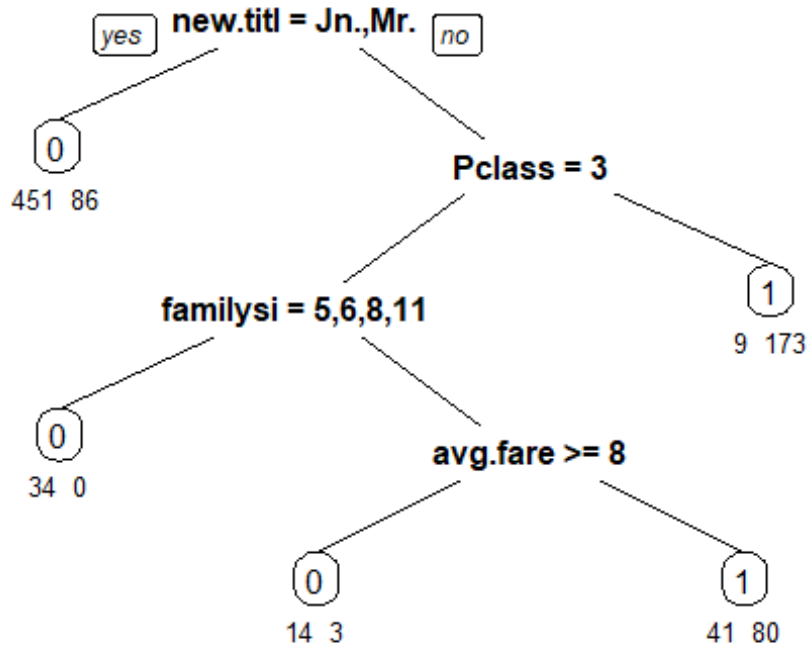
#Run CV and check out results
rpart.3.cv.1 = rpart.cv(94622, rpart.train.3, rf.label, ctrl.3)
rpart.3.cv.1

## CART
##
## 891 samples
## 5 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (3 fold, repeated 10 times)
## Summary of sample sizes: 594, 594, 594, 594, 594, 594, ...
## Resampling results across tuning parameters:
```

```
##
##      cp      Accuracy      Kappa
## 0.00000000 0.8234568 0.6186544
## 0.01593063 0.8318743 0.6359498
## 0.03186126 0.8316498 0.6356703
## 0.04779189 0.8219978 0.6192369
## 0.06372252 0.7970819 0.5703850
## 0.07965316 0.7984287 0.5747427
## 0.09558379 0.7949495 0.5686167
## 0.11151442 0.7940516 0.5669656
## 0.12744505 0.7934905 0.5661919
## 0.14337568 0.7934905 0.5661919
## 0.15930631 0.7934905 0.5661919
## 0.17523694 0.7934905 0.5661919
## 0.19116757 0.7934905 0.5661919
## 0.20709821 0.7934905 0.5661919
## 0.22302884 0.7934905 0.5661919
## 0.23895947 0.7934905 0.5661919
## 0.25489010 0.7934905 0.5661919
## 0.27082073 0.7934905 0.5661919
## 0.28675136 0.7934905 0.5661919
## 0.30268199 0.7934905 0.5661919
## 0.31861262 0.7934905 0.5661919
## 0.33454325 0.7934905 0.5661919
## 0.35047389 0.7934905 0.5661919
## 0.36640452 0.7934905 0.5661919
## 0.38233515 0.7934905 0.5661919
## 0.39826578 0.7934905 0.5661919
## 0.41419641 0.7789001 0.5227001
## 0.43012704 0.7585859 0.4606306
## 0.44605767 0.7261504 0.3597443
## 0.46198830 0.7016835 0.2826611
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.01593063.

#Plot
prp(rpart.3.cv.1$finalModel, type = 0, extra = 1, under = T)

## Warning: Bad 'data' field in model 'call' field.
##      To make this warning go away:
##      Call prp with roundint=FALSE,
##      or rebuild the rpart model with model=TRUE.
```



```
#####
##### Submitting, scoring and some analysis #####
#####

#Rpart

#Subset our test records and features
test.submit.df = data.combined[892:1309, features]

#Make predictions
rpart.3.preds = predict(rpart.3.cv.1$finalModel, test.submit.df, type =
"class")
table(rpart.3.preds)

## rpart.3.preds
##  0  1
## 263 155

#Write out a CSV file for submission to kaggle
submit.df = data.frame(PassengerId = 892:1309, Survived = rpart.3.preds)

write.csv(submit.df, file = "Rpart2.csv", row.names = F)

# random forest
```

```

features = c("Pclass", "new.title", "familysize", "ticket.party.size",
"avg.fare")
rf.train.temp = data.combined[1:891, features]

set.seed(1234)
rf.temp = randomForest(x = rf.train.temp, y = rf.label, ntree = 1000)
rf.temp

##
## Call:
## randomForest(x = rf.train.temp, y = rf.label, ntree = 1000)
##              Type of random forest: classification
##              Number of trees: 1000
## No. of variables tried at each split: 2
##
##              OOB estimate of  error rate: 16.61%
## Confusion matrix:
##      0   1 class.error
## 0 500  49  0.08925319
## 1   99 243  0.28947368

test.submit.df = data.combined[892:1309, features]

# Make predictions
rf.preds = predict(rf.temp, test.submit.df)
table(rf.preds)

## rf.preds
##      0      1
## 277 141

# Write out a CSV file
submit.df = data.frame(PassengerId = 892:1309, Survived = rf.preds)

write.csv(submit.df, file = "RF2.csv", row.names = F)

```